

First we will read the data, noting the different columns and accordingly, label the columns.

1. Year: Year of observation.
2. Month: Month of observation.
3. Day: Day of observation.
4. Fractional Year: Decimal representation of the date.
5. Sunspot Number: Daily sunspot count (can be missing -1).
6. Standard Deviation: Uncertainty in sunspot count.
7. Number of Observations: How many observations contributed to the sunspot count.
8. Definitive or Provisional: 1 for final data, 0 for provisional.

```
import pandas as pd

data = pd.read_csv('raw_data.csv', sep = ";", header = None)

data.columns = ['Year', 'Month', 'Day', 'Fractional Year', 'Sunspot Number',
               'Sunspot Std Dev', 'Num Observations', 'Data Quality']

print(data)
```

```

0      Year  Month  Day  Fractional Year  Sunspot Number  Sunspot Std Dev  \
0      1818     1     1      1818.001          -1          -1.0
1      1818     1     2      1818.004          -1          -1.0
2      1818     1     3      1818.007          -1          -1.0
3      1818     1     4      1818.010          -1          -1.0
4      1818     1     5      1818.012          -1          -1.0
...      ...     ...     ...      ...      ...      ...
75601  2024     12    27      2024.988          258          34.1
75602  2024     12    28      2024.990          252          52.2
75603  2024     12    29      2024.993          234          32.2
75604  2024     12    30      2024.996          218          23.6
75605  2024     12    31      2024.999          179          22.9

      Num Observations  Data Quality
0                    0              1
1                    0              1
2                    0              1
3                    0              1
4                    0              1
...      ...      ...
75601            22              0
75602            26              0
75603            16              0
75604            17              0
75605            12              0

[75606 rows x 8 columns]
```

The year, month and day columns have been combined as the date column, which was then set as the index for the dataset, so it will be easier to run models on the data.

```
required_columns = ['Year', 'Month', 'Day']

data['Date'] = pd.to_datetime(data[required_columns])

data.set_index('Date', inplace=True)

data.drop(['Year', 'Month', 'Day', 'Fractional Year'], axis=1, inplace=True)

print(data)
```

```

Date      Sunspot Number  Sunspot Std Dev  Num Observations  Data Quality
1818-01-01          -1          -1.0              0              1
1818-01-02          -1          -1.0              0              1
1818-01-03          -1          -1.0              0              1
1818-01-04          -1          -1.0              0              1
1818-01-05          -1          -1.0              0              1
...      ...      ...      ...      ...
2024-12-27          258          34.1              22              0
2024-12-28          252          52.2              26              0
2024-12-29          234          32.2              16              0
2024-12-30          218          23.6              17              0
2024-12-31          179          22.9              12              0

[75606 rows x 4 columns]
```

We remove the sunspot Numbers that are -1, and save the remaining rows in a new csv titled "filtered\_sunspot\_data.csv".

```
data_filtered = data[data['Sunspot Number'] != -1]

data_filtered = data_filtered.reset_index()

required_columns = ['Date', 'Sunspot Number', 'Sunspot Std Dev', 'Num Observations', 'Data Quality']
data_filtered = data_filtered[required_columns]

data_filtered.to_csv('filtered_sunspot_data.csv', index=False)

print("Filtered data has been saved to 'filtered_sunspot_data.csv'")

data = pd.read_csv('filtered_sunspot_data.csv', index_col='Date', parse_dates=True)
print(data)
```

Filtered data has been saved to 'filtered\_sunspot\_data.csv'

Date	Sunspot Number	Sunspot Std Dev	Num Observations	Data Quality
1818-01-08	65	10.2	1	1
1818-01-13	37	7.7	1	1
1818-01-17	77	11.1	1	1
1818-01-18	98	12.6	1	1
1818-01-19	105	13.0	1	1
...	...	...	...	...
2024-12-27	258	34.1	22	0
2024-12-28	252	52.2	26	0
2024-12-29	234	32.2	16	0
2024-12-30	218	23.6	17	0
2024-12-31	179	22.9	12	0

[72359 rows x 4 columns]

```
# Select features for anomaly detection
features = data[['Sunspot Number', 'Sunspot Std Dev']]
```

```
from sklearn.ensemble import IsolationForest
# Configure Isolation Forest
iso_forest = IsolationForest(n_estimators=100, contamination=0.05, random_state=42)

# Fit the model
iso_forest.fit(features)
```

IsolationForest

IsolationForest(contamination=0.05, random\_state=42)

```
# Predict anomalies (outliers)
data['Anomaly_Score'] = iso_forest.fit_predict(features)
```

```
# Count anomalies and normal points
anomalies_count = data['Anomaly_Score'].value_counts()
print(anomalies_count)
```

Anomaly\_Score

1	68744
-1	3615

Name: count, dtype: int64

```
# Filter out anomalies
anomalies = data[data['Anomaly_Score'] == -1]
normal_data = data[data['Anomaly_Score'] == 1]

# Preview anomalies
print(anomalies)
```

Anomalies:

Date	Sunspot Number	Sunspot Std Dev	Num Observations	Data Quality
1818-05-29	202	18.0	1	1
1826-11-18	265	20.6	1	1
1826-12-07	268	20.8	1	1
1826-12-09	227	19.1	1	1
1828-03-14	215	18.6	1	1
...	...	...	...	...
2024-12-27	258	34.1	22	0
2024-12-28	252	52.2	26	0
2024-12-29	234	32.2	16	0
2024-12-30	218	23.6	17	0

2024-12-31      179      22.9      12      0

Anomaly\_Score

Date

1818-05-29	-1
1826-11-18	-1
1826-12-07	-1
1826-12-09	-1
1828-03-14	-1
...	...
2024-12-27	-1
2024-12-28	-1
2024-12-29	-1
2024-12-30	-1
2024-12-31	-1

[3615 rows x 5 columns]


```
import matplotlib.pyplot as plt

# Plot sunspot number with anomalies highlighted
plt.figure(figsize=(10,6))

# Plot the normal points
plt.plot(normal_data.index, normal_data['Sunspot Number'], 'b-', label='Normal Data')

# Plot the anomalies
plt.plot(anomalies.index, anomalies['Sunspot Number'], 'ro', label='Anomalies')

plt.title('Sunspot Number with Anomalies Detected by Isolation Forest')
plt.xlabel('Date')
plt.ylabel('Sunspot Number')
plt.legend()
plt.show()
```




```
import matplotlib.pyplot as plt

# Scatter plot for normal and anomaly points
plt.figure(figsize=(10,6))

# Plot normal data points
plt.scatter(normal_data.index, normal_data['Sunspot Number'], label='Normal Data', color='blue')

# Plot anomalies with different marker
plt.scatter(anomalies.index, anomalies['Sunspot Number'], label='Anomalies', color='red', marker='x')

plt.title('Sunspot Number with Anomalies Detected by Isolation Forest')
plt.xlabel('Date')
plt.ylabel('Sunspot Number')
plt.legend()
plt.show()
```



```
# Calculate rolling mean to smooth out the series
rolling_mean = data['Sunspot Number'].rolling(window=30).mean()

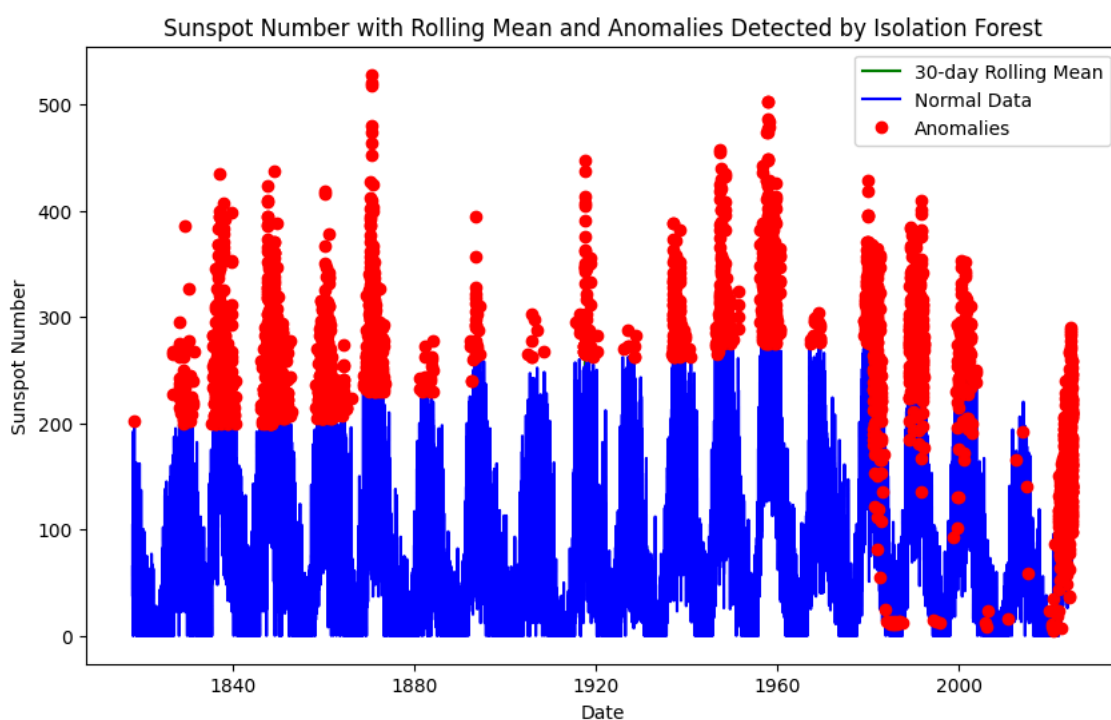
# Plot the rolling mean along with anomalies
plt.figure(figsize=(10,6))

# Plot rolling mean
plt.plot(data.index, rolling_mean, label='30-day Rolling Mean', color='green')

# Plot normal sunspot numbers
plt.plot(normal_data.index, normal_data['Sunspot Number'], 'b-', label='Normal Data')

# Plot anomalies
plt.plot(anomalies.index, anomalies['Sunspot Number'], 'ro', label='Anomalies')

plt.title('Sunspot Number with Rolling Mean and Anomalies Detected by Isolation Forest')
plt.xlabel('Date')
plt.ylabel('Sunspot Number')
plt.legend()
plt.show()
```




```
import seaborn as sns

# KDE plot for normal data and anomalies
plt.figure(figsize=(10,6))

# Plot KDE for normal sunspot numbers
sns.kdeplot(normal_data['Sunspot Number'], label='Normal Data', color='blue')

# Plot KDE for anomalous sunspot numbers
sns.kdeplot(anomalies['Sunspot Number'], label='Anomalies', color='red')

plt.title('KDE Plot of Sunspot Numbers for Normal and Anomalous Data')
plt.xlabel('Sunspot Number')
plt.legend()
plt.show()
```



```
# Calculate the percentage change of sunspot numbers
data['Percentage Change'] = data['Sunspot Number'].pct_change() * 100

# Define a threshold for a "sudden spike" (e.g., change > 50%)
spikes = data[data['Percentage Change'] > 50]

# Filter anomalies that are also spikes
anomalous_spikes = anomalies[anomalies.index.isin(spikes.index)]

# Display sudden spikes that are also anomalies
print(anomalous_spikes)

# Plot these spikes
plt.figure(figsize=(10,6))

# Plot normal data
plt.plot(normal_data.index, normal_data['Sunspot Number'], 'b-', label='Normal Data')


# Highlight anomalies that are spikes
plt.plot(anomalous_spikes.index, anomalous_spikes['Sunspot Number'], 'ro', label='Anomalous Spikes')

plt.title('Anomalous Sudden Spikes in Sunspot Numbers')
plt.xlabel('Date')
plt.ylabel('Sunspot Number')
plt.legend()
plt.show()
```

Date	Sunspot Number	Sunspot Std Dev	Num Observations	Data Quality \
1826-11-18	265	20.6	1	1
1826-12-07	268	20.8	1	1
1828-06-19	242	19.7	1	1
1829-04-21	225	19.0	1	1
1829-06-29	385	24.9	1	1
...	...	...	...	...
2022-06-13	94	19.9	57	1
2022-10-21	57	18.4	34	1
2023-11-21	153	18.9	35	1
2024-02-23	104	27.4	34	1
2024-03-18	120	17.4	49	1

Date	Anomaly_Score
1826-11-18	-1
1826-12-07	-1
1828-06-19	-1
1829-04-21	-1
1829-06-29	-1
...	...
2022-06-13	-1
2022-10-21	-1
2023-11-21	-1
2024-02-23	-1
2024-03-18	-1

[174 rows x 5 columns]



```
# Calculate the 10th percentile to define unusually low values
low_threshold = data['Sunspot Number'].quantile(0.20)

# Filter anomalies that have unusually low values
unusually_low_anomalies = anomalies[anomalies['Sunspot Number'] < low_threshold]

# Display these unusually low anomalies
print(unusually_low_anomalies)

# Plot these anomalies
plt.figure(figsize=(10,6))

# Plot normal data
plt.plot(normal_data.index, normal_data['Sunspot Number'], 'b-', label='Normal Data')

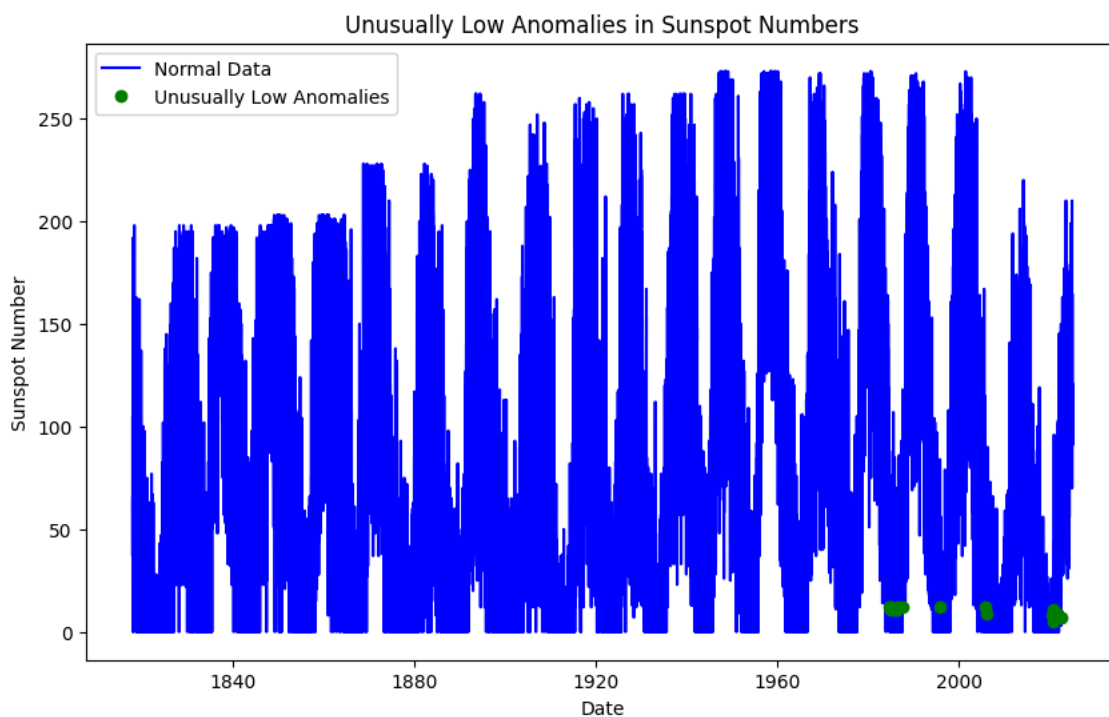
# Highlight unusually low anomalies
plt.plot(unusually_low_anomalies.index, unusually_low_anomalies['Sunspot Number'], 'go', label='Unusually Low Anomalies')

plt.title('Unusually Low Anomalies in Sunspot Numbers')
plt.xlabel('Date')
plt.ylabel('Sunspot Number')
plt.legend()
plt.show()
```



Date	Sunspot Number	Sunspot Std Dev	Num Observations	Data Quality \
1984-09-15	12	10.9	8	1
1984-12-31	11	10.6	11	1
1985-01-08	12	11.4	16	1
1985-08-17	11	10.0	19	1
1985-12-08	11	10.3	10	1
1986-04-01	11	10.2	18	1
1986-10-16	12	11.6	15	1
1987-06-11	12	10.0	24	1
1995-11-03	12	10.8	18	1
2005-10-23	12	10.9	19	1
2006-02-11	9	9.4	17	1
2020-07-07	8	10.4	38	1
2020-08-03	11	11.6	45	1
2020-11-01	5	8.7	33	1
2021-01-28	7	9.0	32	1
2021-08-01	8	10.2	33	1
2021-08-13	9	11.1	45	1
2022-06-07	7	11.0	31	1

Date	Anomaly_Score
1984-09-15	-1
1984-12-31	-1
1985-01-08	-1
1985-08-17	-1
1985-12-08	-1
1986-04-01	-1
1986-10-16	-1
1987-06-11	-1
1995-11-03	-1
2005-10-23	-1
2006-02-11	-1
2020-07-07	-1
2020-08-03	-1
2020-11-01	-1
2021-01-28	-1
2021-08-01	-1
2021-08-13	-1
2022-06-07	-1



```
# Get the time differences between consecutive anomalies
anomalies['Time Diff'] = anomalies.index.to_series().diff().dt.days

# Plot histogram of time gaps between anomalies
plt.figure(figsize=(10,6))
plt.hist(anomalies['Time Diff'].dropna(), bins=20, color='purple', alpha=0.7)

plt.title('Distribution of Time Gaps Between Anomalies')
plt.xlabel('Days Between Anomalies')
plt.ylabel('Frequency')
plt.show()
```

```
<ipython-input-15-d2b68ffe0f60>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-anomalies\['Time Diff'\] = anomalies.index.to\\_series\(\).diff\(\).dt.days](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-anomalies['Time Diff'] = anomalies.index.to_series().diff().dt.days)

